

Venture Capital Backing and Its Relationship to Innovation Quality, Scope, and Science Intensity

Oscar Seltzer^{a,1}

This manuscript was compiled on August 30, 2026

Venture capital (VC) is credited with accelerating high-risk, high-reward technological breakthroughs, yet empirical evidence tracing how private backing shapes these innovations via the impact, legal, and scientific substance of intellectual property remains incomplete. From 7.8 million patents, we constructed a set of 500,000 U.S. patents granted between 2000 and 2020, matched to venture-financing events and disambiguated USPTO assignee, citation, and claim-level records. We evaluated patents across three dimensions: peer-normalized forward citation rate, claim-based legal defensibility, and non-patent-literature (NPL) science intensity. *T*-tests indicate that VC-backed patents receive more than double the peer-normalized forward citations of non-VC patents (2.03 vs. 0.98, $p = 0.014$), carry broader claim sets (22.2 vs. 16.0 claims, $p < 10^{-185}$), and cite academic literature at a higher rate (44.6% vs. 36.9% NPL ratio, $p = 0.0047$). However, the VC benefit is not uniform. Rather, VC backing generates high citation multiples in frontier fields such as DeepTech, shown in CPC sections Electricity (2.68 $\times$) and Physics & Computing (2.34 $\times$). Meanwhile, other less-frontier sectors, such as mechanical engineering and operational transport applications, exhibit no citation advantage over incumbent peers. Further, while VC-backed companies have consistently had more claims than non-VC peers, following policy reform with the 2011 America Invents Act, this difference doubled in only three years, demonstrating increased adaptability to procedural reforms.

venture capital | deep tech | patent citations | patent scope | innovation policy

1. Introduction and Related Work

VC has a relatively unique position among both money-lenders and the modern national innovation system. Unlike commercial lenders that require collateral and continuous, uniform repayments, or public companies that work on fixed quarterly earning cycles, early-state venture works within an extreme power-law return regime. A standard notion within venture suggest that out of thirty portfolio investments in a fund, more than half will fail or return nominal capital, a fraction will generate single-digit multiples, and a single outlier will account for the majority of the fund's total financial return.

This reliance on positive outliers creates strong incentives for VCs. VCs can't optimize their portfolio through small refinements. Instead, the most economical—and many times the only—option is to identify and underwrite frontier technology that are establishing or could establish a dominant market position. Consequently, venture-backed companies are encouraged by their board members to secure broad, commercially defensible IP moats with the goal of both protecting the company from competitors and helping to facilitate an eventual M&A exit.

Significance

Venture capital (VC) funding is behind a large portion of major recent companies (Apple, Google, Anthropic, SpaceX, etc.) most of which have pushed technology forward in large ways, leading to the thought that VC backing may have a relationship to innovation. Analyzing a sample cohort of 500,000 U.S. patents (2000–2020), we show that VC-backed patents receive double the peer-normalized forward citations, carry broader legal claims, and draw more on academic science than non-VC patents. Further, these relationships are non-uniform: citation gains are concentrated in DeepTech, which contains CPC Sections Physics and Computing, and Electricity, while other sectors such as LogTech and HardTech (CPC sections Operations and Transport, and Mechanical Engineering) show no advantage. These findings demonstrate that VC's impact is sector-specific, which can provide empirical guidance for federal innovation policy, patent system reform, and institutional capital allocation.

Author affiliations: ^aQuantitative Social Science Program, Dartmouth College, Hanover, NH 03755

O.S. conceived the study, designed the empirical methodology, built the data pipeline, performed statistical analyses, generated all visualizations, and wrote the manuscript. H.C. advised the project providing insights on potentially relevant figures.

The author declares no competing interests.

¹To whom correspondence should be addressed. E-mail: oscar.seltzer.28@dartmouth.edu

Many economics papers have examined the relationship between VC and patenting behavior. In their first contribution to this matter, Kortum and Lerner (1) established that venture funding is associated with a disproportionate increase in patenting volume per dollar of R&D expenditure relative to corporate R&D spend. This suggests that VC provides a degree of governance and screening mechanism beyond purely providing capital. Further papers have expanded on this looking into managerial and operational mechanisms. Bernstein et al. (2) demonstrated that hands-on VC monitoring and most significantly on-site involvement significantly increases patent volume and citation quality. Fortner, Kortum and Lerner (3) also showed that multi-decade patent surges reflect not only legal reforms, but also research management and investment strategies.

However, at the same time, the literature mentions that patent volume can not be equated to technological disruption. Hall and Ziedonis (4) demonstrated the 'patent paradox,' which shows how firms in capital-intensive sectors patent more often purely for defensive leverage rather than direct technological development. Theoretical models by Akcigit and Kerr (5) and findings by Wu et al. (6) show that legacy incumbents focus on internal, incremental development whereas new entrants focus on disruptive innovation. Hall et al. (7) and Lerner and Seru (8) both stress the importance of correcting for classification biases, vintage, and truncation when working within the patent citation data—something that was attempted to minimize in this project through taking a more recent, random sample from the entire database. Finally, in cyclical downturns, Howell et al. (9) established that venture-backed innovation has increased levels of resilience.

While prior literature have answered many questions about the relationship of venture-backing and innovation via patents, many core questions regarding the substantive characteristic of venture-backed IP remain. Most prior studies focused on only one element of patent counts, forward citations, or legal claims without looking at them together. Still, the largest difference is that most prior literature looks at venture-backing as a single variable instead of comparing and contrasting venture-backing in different technological sectors.

This paper address these gaps by evaluating a sample of 500,000 U.S. patents granted between 2000 and 2020 taken from a broader database of 7.8 million patents from 1976-2026. USPTO bulk database records were linked with venture investment events and then analyzed through three dimensions of patent output: peer-normalized forward citation impact, claim-based legal scope, and science intensity. A particular emphasis was given to looking at the differences between sectors across Cooperative Patent Classification (CPC) technology sections and analyzing the impact of policy shocks when examining the 2011 Leahy-Smith America Invents Act (AIA).

Our findings demonstrate that VC's technological benefits are sector-specific. While venture-backed patents tend to exhibit advantages in citation impact, legal scope, and scientific grounding across the aggregate sample, this advantage is concentrated in DeepTech, hardware, semiconductor, computing, and other similar domains. On the other hand, venture-backed patents perform on par with or below non-venture peers in mature engineering and transportation sectors. Finally, venture-backed patents experienced a massive uptick in claims following AIA. These results help to provide guidance for capital allocators and national innovation policymakers that seek to optimize innovation, growth, and success of American companies across diverse sectors.

2. Data

The process of constructing the initial complete dataset (i.e. prior to sampling), which contained data capable of tracking disambiguated assignee entities, forward and backward citation linking for every patent, and claim characteristics, was one of the more arduous parts of this project. The primary files released by the USPTO and PatentsView contain hundreds of millions of relational records, many of which need to be cross-checked to create citation links. However, if one were to expand on this project, they could go down the route of a USPTO API. Unfortunately, we were unable to get approved for said integration during the time horizon of this project.

The analytical panel consists of a sample of 500,000 U.S. patents granted between 2000 and 2020, constructed from three primary empirical data sources:

- **PatentsView Bulk Releases:** Relational bulk tables including patent classifications (`g_cpc_current`), forward and backward patent citations (`g_us_patent_citation`, encompassing over 130 million citation dyads), non-patent literature references (`g_other_reference`), and disambiguated corporate assignees (`g_assignee_disambiguated`).
- **PatentVC Enhanced Linkage:** A publicly curated venture panel (`patentvc_enhanced_4var.dta`) matching PitchBook and VentureSource institutional financing rounds to corresponding USPTO patent grant numbers.
- **USPTO Patent Claims Statistics:** Historical claim-level data (`patent_claims_stats.dta.zip`) tracking total claims, independent claim counts, and claim word lengths.

Each patent was characterized by its grant year, application year, primary CPC technology section (A–H, restricted to Human Necessities, Operations & Transport, Chemistry & Metallurgy, Mechanical Engineering, Physics & Computing, and Electricity; sections D, Textiles, and E, Fixed Constructions, are excluded for insufficient VC-cohort observations), 4-character CPC subclass, a binary VC-backing indicator, raw and normalized forward citation counts, backward patent citations, NPL citations, total claim counts, the logarithmic transformation of total claim counts, and disambiguated assignee identifiers.

To minimize truncation artifacts that are caused by patent vintage (i.e. older patents tend to have more forward citations) and systemic differences in citation velocity across sectors (7, 8), we construct a **peer-normalized forward citation index**:

$$\text{Normalized Citations}_i = \frac{\text{Raw Forward Citations}_i}{\mathbb{E}[\text{Raw Forward Citations} \mid \text{CPC Subclass}_i, \text{Grant Year}_i]} \quad [1]$$

Under this formula, a value of 1.0 represents average citation performance for patents of the same technology class and grant year. However, due to a substantive difference in the percentage of VC-backed patents to non-VC-backed patents, non-VC-backed patents normalize to a value of ~ 1.0 .

Science intensity is captured via the Non-Patent Literature (NPL) ratio:

$$\text{NPL Ratio}_i = \frac{\text{NPL Citations}_i}{\text{Backward Patent Citations}_i + \text{NPL Citations}_i} \quad [2]$$

This bounded metric $[0, 1]$ reflects the proportion of preceding technology grounded in academic literature, conference proceedings, and scientific papers rather than existing patent disclosures.

Finally, legal defensibility is captured through total claim counts and the logarithmic transformation $\ln(\text{Total Claims} + 1)$; this logarithmic transformation is used to correct for a high degree of positive skewness.

3. Methods

The ETL pipeline was created to prevent crashes due to file size given the scale of the underlying USPTO bulk files and limited compute capabilities. Processing proceeds in three streaming stages using DuckDB’s vectorized OLAP engine over uncompressed local SSD storage and PyArrow chunked conversions for Stata (`.dta`) files:

Step 1: DuckDB streams raw CPC, citation, and NPL reference tables to construct patent-level forward, backward, and NPL citation totals for a 500,000 patent sample from 2000–2020, exporting `step1.metrics.parquet`.

Step 2: Patent claims statistics are streamed from the Stata claims file in 5-million-row chunks via PyArrow, matched to the cohort by unique patent identifier, and exported as `step2.claims.parquet`.

Step 3: Step 1 and Step 2 outputs are merged with the PatentVC linkage panel and disambiguated assignee data; forward citations are benchmarked against CPC-subclass $\times$ grant-year peer means to generate `normalized_forward_citations`, and the panel is exported as `cleaned_patent_panel.parquet / .csv`.

On the assembled dataset, we executed two-sample t -tests comparing group means for normalized forward citations, total claims, and NPL ratios. Rather than a Student’s t -test, a Welch’s t -test was used as it does not assume equal variances, which was necessary due to the sample size imbalance between the VC-backed cohort ($n = 9,536$ total patents; $n = 148$ non-zero citation records) and the non-VC baseline ($n = 490,464$ total patents; $n = 7,864$ citation records).

Additionally, we estimated the empirical kernel density distributions of science intensity and log-claims by VC status, disaggregated mean normalized citations by CPC technology section against a 1.0 parity benchmark, and examined the number of claims per patent from 2008–2014 around the enactment of the 2011 America Invents Act (AIA).

4. Results: Writing

4.1. Aggregate Comparisons. Across the 500,000-patent sample, VC-backed patents demonstrate statistically significant premiums across all three primary outcome dimensions. First, VC-backed patents receive more than double the peer-normalized forward citations of non-VC patents (2.026 vs. 0.981, $t = 2.49$, $p = 0.0140$). Since citations are standardized against 4-character CPC subclass $\times$ grant-year peer groups, this is a true difference that is not merely explained by faster-moving technology categories. Second, VC-backed patents have broader claim sets (22.23 vs. 16.03 claims, $t = 29.71$, $p < 10^{-185}$). In other words, VC-backed patents are more defensively drafted and likely, as a result, more resource-intensive to craft. Third, VC-backed patents have a higher NPL ratio (0.446 vs. 0.369, $t = 2.87$, $p = 0.0047$), perhaps indicating that venture-funded technologies build more on scientific discoveries.

4.2. Sectoral Heterogeneity. Disaggregating citation impact across CPC technology sections shows that the venture-backing premium is sector-specific. Instead of VC-backed patents exhibiting a uniform citation advantage, they are heavily concentrated in sectors such as DeepTech, hardware, semiconductor, and computational fields. In Section H (Electricity, encompassing semiconductors, microelectronics, and energy storage) and Section G (Physics & Computing, encompassing software algorithms and quantum architectures), VC-backed patents achieve $2.68\times$ and $2.34\times$ their peer-group citation benchmarks, respectively. There are moderate premiums in Section C (Chemistry & Metallurgy, $1.54\times$) and Section A (Human Necessities/Biotechnology, $1.48\times$). However, VC-backed inventions in traditional industrial categories, such as Section B (Operations & Transport, $0.77\times$) and Section F (Mechanical Engineering, $0.76\times$), underperform the 1.0 benchmark.

4.3. Distributions and Policy Responsiveness. KDE of the NPL ratio demonstrates that while non-VC patents cluster near zero, VC-backed patents have a relatively uniform density across the 0.2–0.9 range, which backs up the idea of sustained proximity to open scientific research. While likely a product of defensive patenting strategies, analysis of claim counts shows an upward shift across all quartiles for VC patents, with a median of 19 claims (IQR: 11–29) compared to 14 claims (IQR: 8–21) for non-VC patents. When examining claim drafting around policy reform (the 2011 America Invents Act), another divergence can be seen. Non-VC claim counts remained fairly flat (14–15 claims) throughout the entire 2008–2014 window. Meanwhile, VC-backed claims sank to 15.6 claims in the 2011 reform year before rising sharply to 20.5 in 2012 and eventually reaching a high of 23.4 claims by 2014, suggesting that VC-backed companies are highly sensitive to policy reform, unlike their non-VC counterparts.

5. Results: Figures and Tables

Table 1. Summary statistics for primary outcome variables across the 500,000-patent panel (2000–2020).

Variable	<i>n</i>	Mean	SD	Median	IQR
<i>Norm. fwd. citations</i>	8,012	1.000	2.564	0.231	0.00–1.12
<i>NPL ratio</i>	7,871	0.370	0.352	0.286	0.04–0.67
<i>Total claims</i>	500,000	16.150	14.081	14.000	8.00–21.00
<i>Log claims</i>	500,000	2.624	0.693	2.708	2.19–3.09

Normalized forward citations and NPL ratios reflect patents with complete citation metadata and non-zero backward-reference denominators.

Table 2. Welch’s two-sample *t*-test comparisons between VC-backed and non-VC patents.

Outcome Metric	VC Mean (<i>n</i>)	Non-VC Mean (<i>n</i>)	<i>t</i> -statistic	<i>p</i> -value
<i>Norm. fwd. citations</i>	2.026 (148)	0.981 (7,864)	2.49	0.0140
<i>Total claims</i>	22.230 (9,536)	16.032 (490,464)	29.71	< 10 ^{−180}
<i>Log claims</i>	2.894 (9,536)	2.619 (490,464)	37.45	< 10 ^{−180}
<i>NPL ratio</i>	0.446 (146)	0.369 (7,725)	2.87	0.0047

Degrees of freedom reflect non-missing data subsets under Welch’s unequal variance formulation.

Figure 1: (Horizontal Coefficient) Mean Normalized Forward Citation Impact Across CPC Technology Sections




Figure 1: Mean normalized forward citations for VC-backed patents by CPC section (2000–2020). Normalized citations are computed relative to the peer benchmark (CPC subclass × grant year mean = 1.0, dashed line). Markers indicate point estimates with standard error whiskers. Non-VC patents are omitted as their normalized mean sits at approximately 1.0 by definition. Sections D (Textiles) and E (Fixed Constructions) are excluded due to insufficient observation counts in the VC cohort sample.

Fig. 1. Mean normalized forward citations for VC-backed patents by CPC technology section, benchmarked against peer parity (dashed line at 1.0; subclass × grant-year mean). Error bars indicate standard errors. Sections D (Textiles) and E (Fixed Constructions) are excluded due to insufficient VC-cohort observations.

6. Discussion

The findings presented here provide three primary insights into the relationship between venture capital and patent output:

First, the DeepTech citation multiplier documented in Figure 1 demonstrates that venture capital’s contribution to technological progress is domain-dependent. In Electricity (2.68×) and Physics & Computing (2.34×), venture backing, milestone funding, and technical syndicates work exceptionally well. In frontier sectors, where market leadership is driven by novel materials, new physics, etc., VC is a perfect fit to provide the needed risk-tolerant equity to fund protracted development cycles. These cycles lead to highly cited patents that serve as prior art for subsequent innovation.




figure 2a: Empirical distribution of science intensity comparing vc-backed and non-vc patents (2000–2020). science intensity is defined as the ratio of non-patent literature citations to total backward citations.

Fig. 2. Empirical kernel density estimates of science intensity (NPL ratio) comparing VC-backed and non-VC patents (2000–2020). Non-VC patents concentrate heavily near zero, while VC patents exhibit sustained density across high-NPL values.




Figure 4: Distribution of patent total claims count across cohorts (2000–2020). Boxplots display medians and interquartile ranges (outliers suppressed for visual clarity).

Fig. 3. Boxplot distribution of total patent claim counts by funding status (2000–2020; extreme outliers suppressed for visual clarity). The VC-backed median (19) and IQR (11–29) substantially exceed the non-VC median (14) and IQR (8–21).




Figure 3: Evolution of patent claim counts during the years preceding and following the 2011 America Invents Act (AIA) (2008–2014). Total claims serve as a proxy for legal scope and complexity. The dashed vertical line marks the enactment of the 2011 reform.

Fig. 4. Mean total patent claim counts by grant year and funding status (2008–2014) spanning the enactment of the 2011 America Invents Act (dashed line). VC-backed claims display a sharp compression at the reform followed by an aggressive expansion to series highs.

Conversely, in more mature engineering sectors (Sections B and F), where progress is more reliant on supply chains and manufacturing optimizations, VC has no advantage over a corporate incumbent.

Second, the elevated claim counts (Figure 3) and higher science intensity (Figure 2) demonstrate a specific legal-commercial strategy. Venture-backed startups construct broad claim trees to maximize defensive moats against incumbent litigation. These claims require anchoring in peer-reviewed scientific literature to provide increased protection against prior-art invalidation during post-grant review and infringement proceedings.

Third, the claim dynamics observed around the 2011 AIA (Figure 4) demonstrate that venture-backed companies are highly sensitive and more adaptive to patent reforms. While non-VC filers had no change in claim volume, VC-backed firms adapted quickly to the first-inventor-to-file regime by expanding claim counts to lock in more defensible IP.

All of this combined suggests a few things: (1) VC is critical for frontier innovation, but less so for other sectors, and (2) patent policy reform such as the 2011 AIA may actually be detrimental to innovation. Instead of helping general inventors, it seems to have only increased defensibility and therefore limited innovation through competition.

7. Limitations and Future Work

Still, there are several caveats that limit the interpretation of these results. First, while the sample pulled encompasses 500,000 patents, the venture-backed subsample is relatively small ($n = 146$), which leads to wide confidence intervals when looking at specific sectors. Second, the matching process was done via PatentVC's curated sample instead of a complete census, which may induce selection effects from firm size and vintage as their exact data collection processes are unknown. Third, all comparisons remain associational rather than causal, as venture firms search for and select high-caliber teams who might produce superior IP regardless of being funded.

Future research should expand the examined data to increase matched citation sample sizes, integrate round-level funding stages—through using an SEC API—and look at quasi-experimental policy shocks (such as state-level R&D tax credits or SBIR grant cutoffs) to prove if the relationships are causal or not. Additionally, decomposing claim counts into independent versus dependent claims and controlling for assignee firm age would further clarify if the VC scope premium reflects venture backing or the dynamics of first-time patenting entities.

8. Code and Repository Availability

All ETL extraction pipelines, econometric analysis scripts, and visualization code are open-source and publicly available on GitHub at https://github.com/Oscar-Seltzer/qss20_26x_vc. A web companion presenting the figures and underlying datasets is deployed at <https://vc-innovation.vercel.app/>. Raw bulk data files from PatentsView and the USPTO are accessible through the PatentsView public download portal.

9. Agentic Analysis

<https://share.gemini.google/vtvz35Sh7DCd>

Note: This is merely one of my more recent AI-transcripts that reflects the general prompts given. This specific transcript is an amended version of a prior transcript with spelling mistakes fixed to aid in legibility.

I tried to avoid ever asking the assistant to one-shot any code. In general, my coding pipeline was to: (1) come up with an idea of what I wanted, (2) discuss that idea with an AI for it to give pros and cons and options, (3) create the code scaffolding (normally through comments), (4) have the AI fill out the scaffolding, (5+) iterations of modifying sections of the code until I was happy with it all.

Things tended to go wrong when I would instead ask the assistant to one-shot something (i.e. I gave my initial idea to it and told it to code out everything). When that happened it would tend to focus on things I didn't want it to (in the case of making figures), code in an overly-obtuse manner, create many bugs, or just do something completely different from what I asked.

This was most prominent where in the case of building out the website after giving the initial prompt I got back something completely different from what I prompted. After that, I went into each core file and worked on it with the assistant piece by piece before combining everything back together into the website.

In general, I also tried to avoid blindly accepting any changes. When my assistant suggested a change I would counter with the reasoning, pros, and cons for said change. That said, when it came to the process of initially unzipping, reading, and processing the data I relied more heavily on the assistant after a few hours and a couple completely different techniques failing. At that point, I had my assistant give me a list of 10 different options and then one-shotted each one. After I did that, I examined which option worked best, had my assistant give me a lesson in how it worked, and then went back to my traditional assisted coding method.

AI assistants were also used in non-coding matters. For example, in the creation of the website, as mentioned earlier, and in doing a grammar and flow check of content on the website.

1. S Kortum, J Lerner, Assessing the contribution of venture capital to innovation. *The RAND J. Econ.* **31**, 674–692 (2000).
2. S Bernstein, X Giroud, RR Townsend, The impact of venture capital monitoring. *The J. Finance* **71**, 1591–1622 (2016).
3. S Kortum, J Lerner, What is behind the recent surge in patenting? *Res. Policy* **28**, 1–22 (1999).
4. BH Hall, RH Ziedonis, The patent paradox revisited: An empirical study of patenting in the u.s. semiconductor industry, 1979–1995. *The RAND J. Econ.* **32**, 101–128 (2001).
5. U Akcigit, WR Kerr, Growth through heterogeneous innovations. *J. Polit. Econ.* **126**, 1374–1443 (2018).
6. L Wu, D Wang, JA Evans, Large teams develop and small teams disrupt science and technology. *Nature* **566**, 378–382 (2019).
7. BH Hall, AB Jaffe, M Trajtenberg, The nber patent citation data file: Lessons, insights and methodological tools, (National Bureau of Economic Research), Working Paper 8498 (2001).
8. J Lerner, A Seru, The use and misuse of patent data: Issues for finance and beyond. *The Rev. Financial Stud.* **35**, 2667–2704 (2022).
9. ST Howell, J Lerner, R Nanda, RR Townsend, How resilient is venture-backed innovation? evidence from four decades of u.s. patenting, (National Bureau of Economic Research), Working Paper 27150 (2020).